

Concurrent programming


Introduction

Companion slides for
The Art of Multiprocessor Programming
by Maurice Herlihy, Nir Shavit, Victor Luchangco,
and Michael Spear

Modified by Piotr Witkowski

Moore's Law

42 Years of Microprocessor Trend Data




Transistor
count still
rising


Clock
speed
flattening
sharply

Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2017 by K. Rupp

Extinct: the Uniprocessor

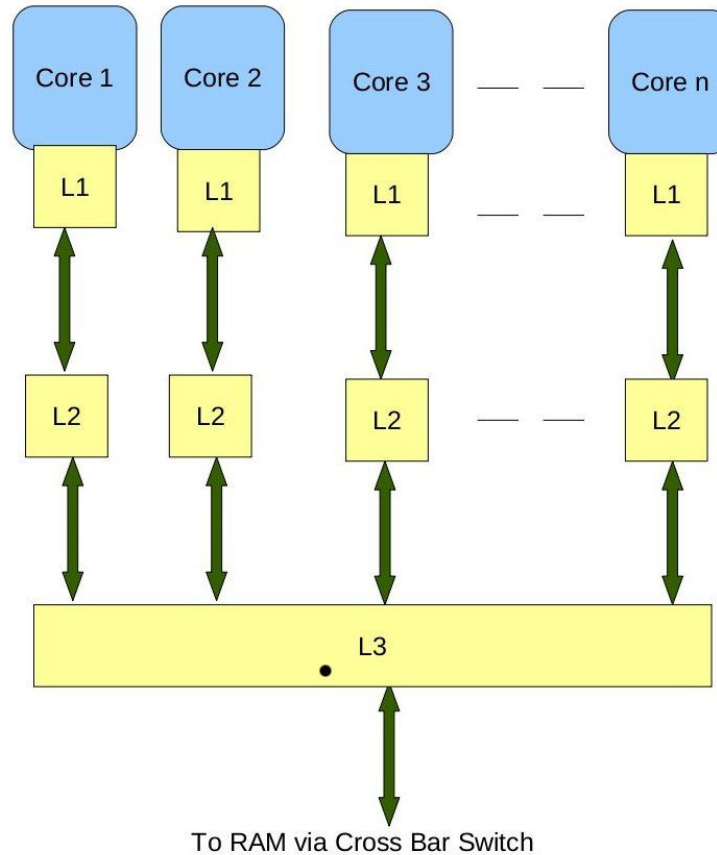


Extinct: The Shared Memory Multiprocessor (SMP)



The New Boss: The Multicore Processor

**All on the
same chip /
chiplet**




Intel Core
i9 upto 18
cores

AMD
Ryzen
upto 64
cores


Why do we care?

- Time no longer cures software bloat
 - The “free ride” is over
- When you double your program’s path length
 - You can’t just wait 6 months
 - Your software must somehow exploit twice as much concurrency

Traditional Scaling Process

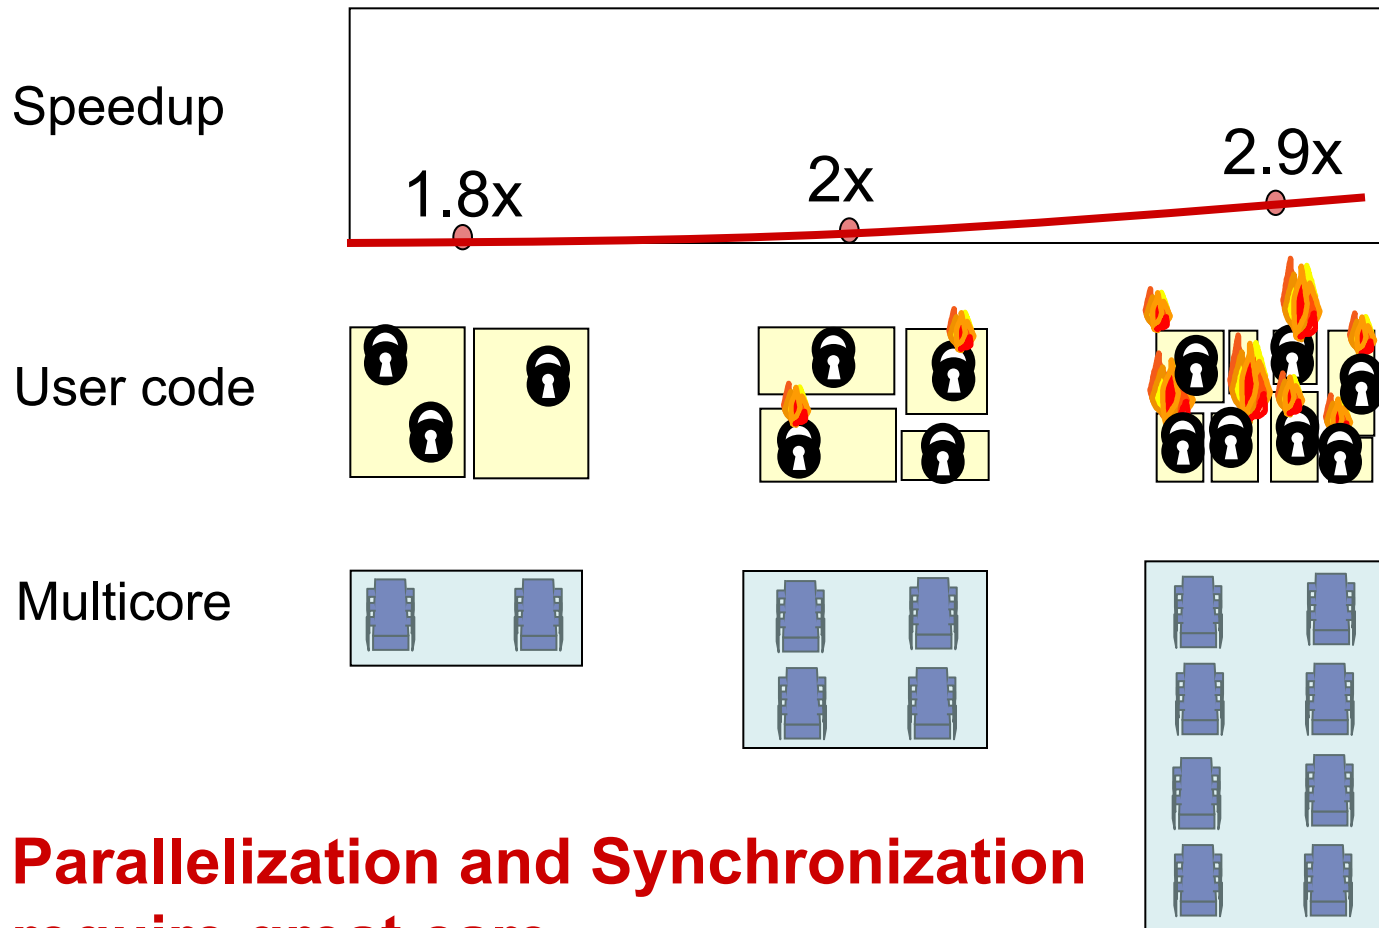


Ideal Scaling Process



Unfortunately, not so simple...

Actual Scaling Process




**Parallelization and Synchronization
require great care...**


Concurrent Programming: Course Overview

- Fundamentals
 - Models, algorithms, impossibility
- Real-World programming
 - Architectures
 - Techniques


Sequential Computation



Concurrent Computation



Asynchrony



- Sudden unpredictable delays
 - Cache misses (*short*)
 - Page faults (*long*)
 - Scheduling quantum used up (*really long*)

Model Summary

- Multiple *threads*
 - Sometimes called *processes*
- Single shared *memory*
- *Objects* live in memory
- Unpredictable asynchronous delays

Road Map

- We are going to focus on principles first, then practice
 - Start with idealized models
 - Look at simplistic problems
 - Emphasize correctness over pragmatism
 - “Correctness may be theoretical, but incorrectness has practical impact”


Concurrency Jargon

- Hardware
 - processors \approx cores
- Software
 - threads \approx processes
- Programing
 - concurrent \approx parallel \approx multiprocessor \approx multicore
- Sometimes OK to confuse them, sometimes not.

Parallel Primality Testing

- Challenge
 - Print primes from 1 to 10^{10}
- Given
 - Ten-processor multiprocessor
 - One thread per processor
- Goal
 - Get ten-fold speedup (or close)

Load Balancing



- Split the work evenly
- Each thread tests range of 10^9

Procedure for Thread i

```
void primePrint {  
    int i = ThreadID.get(); // IDs in {0..9}  
    for (j = i*109+1, j<(i+1)*109; j++) {  
        if (isPrime(j))  
            print(j);  
    }  
}
```

Issues


- Higher ranges have fewer primes
- Yet larger numbers harder to test
- Thread workloads
 - Uneven
 - Hard to predict

Issues

- Higher ranges have fewer primes
- Yet larger numbers harder to test
- Thread workloads
 - Uneven
 - Hard to predict
- Need *dynamic* load balancing



Shared Counter



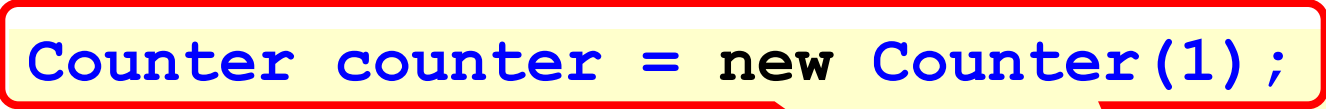
Procedure for Thread i

```
int counter = new Counter(1);

void primePrint {
    long j = 0;
    while (j < 1010) {
        j = counter.getAndIncrement();
        if (isPrime(j))
            print(j);
    }
}
```

Procedure for Thread *i*

```
Counter counter = new Counter(1);
```



```
void primePrint {  
    long j = 0;  
    while (j < 1010) {  
        j = counter.getAndIncrement();  
        if (isPrime(j))  
            print(j);  
    }  
}
```


**Shared counter
object**

Where Things Reside

```
void primePrint (  
    int i = ThreadID.get();  
    // IDs in [0..9]  
    for (j = i*10+1;  
         j<(i+1)*10; j++) {  
        if (isPrime(j))  
            print(j);  
    }  
}
```

code

Local
variables



Procedure for Thread *i*

```
Counter counter = new Counter(1);
```

```
void primePrint {
```

```
    long j = 0;
```

```
    while (j < 1010) {
```

```
        j = counter.getAndIncrement();
```

```
        if (isPrime(j))
```

```
            print(j);
```

```
    }
```

```
}
```

**Stop when every
value taken**

Procedure for Thread i

```
Counter counter = new Counter(1);
```

```
void primePrint {
```

```
    long j = 0;
```

```
    while (j < 1010) {
```

```
        j = counter.getAndIncrement();
```

```
        if (isPrime(j))
```

```
            print(j);
```

```
    }
```

```
}
```

**Increment & return each
new value**

Counter Implementation

```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        return value++;  
    }  
}
```

Counter Implementation

```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        return value++;  
    }  
}
```

**OK for single thread,
not for concurrent threads**

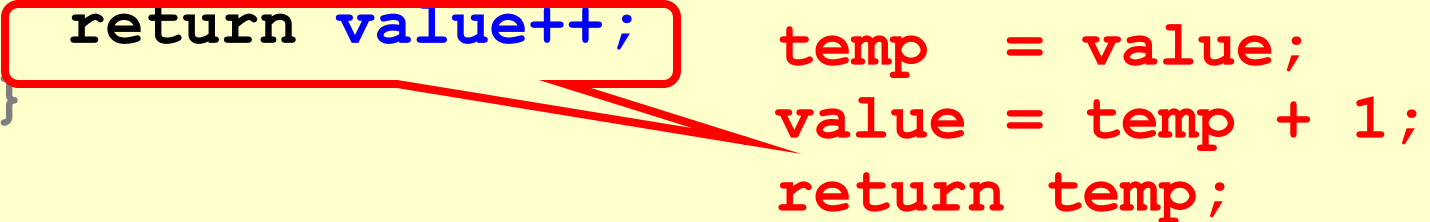
What It Means

```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        return value++;  
    }  
}
```

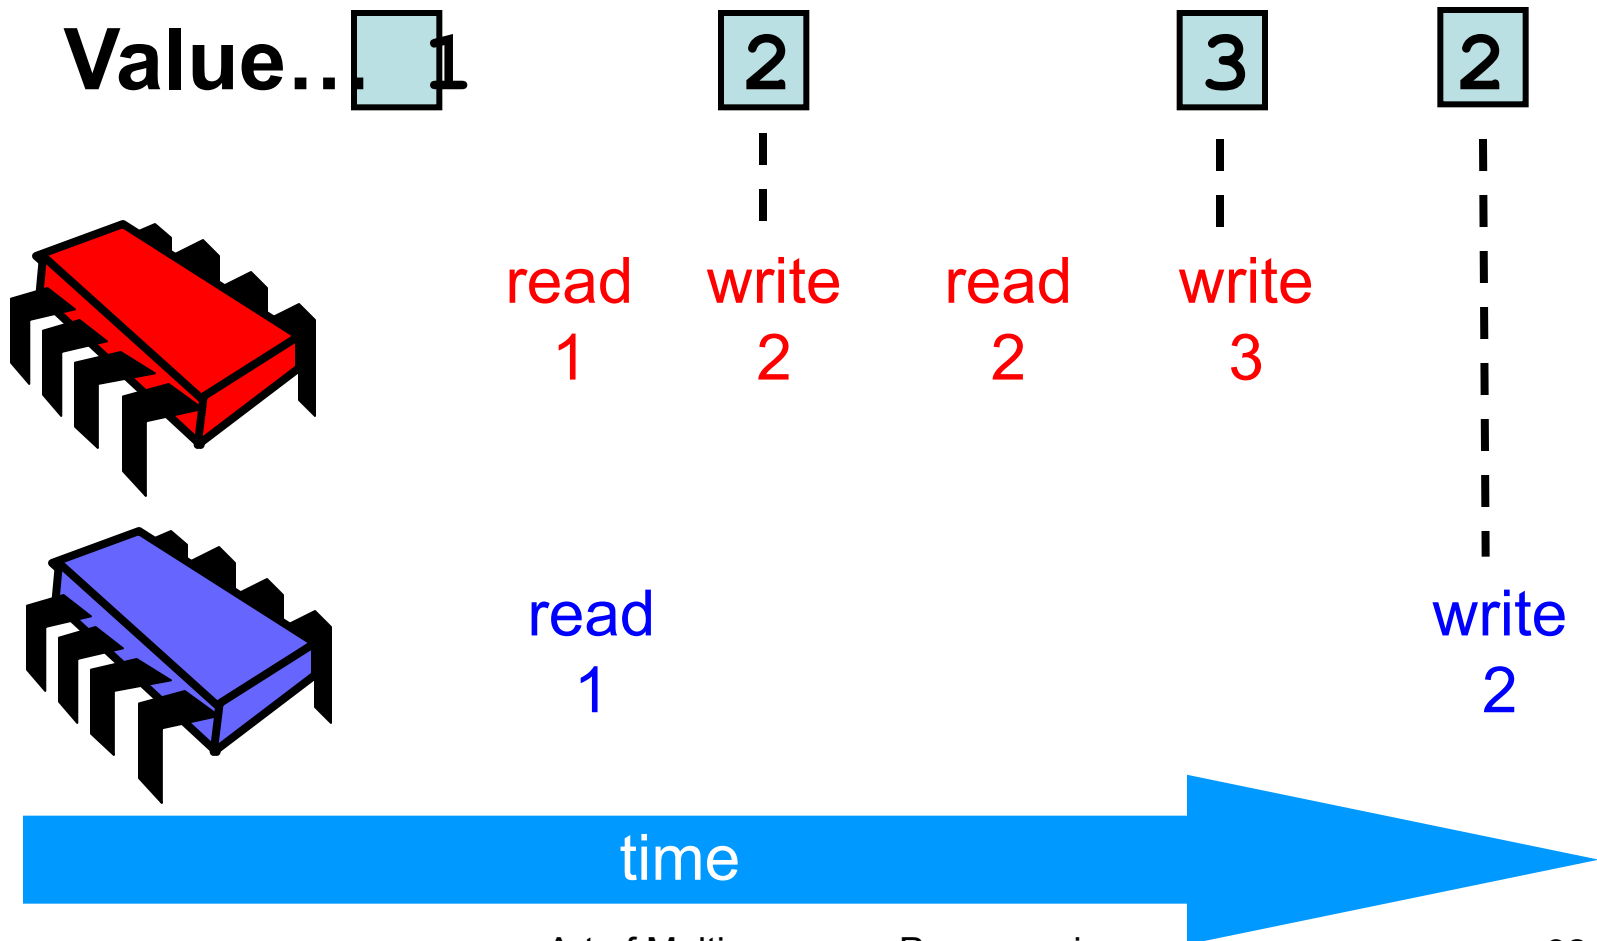
What It Means

```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        return value++;  
    }  
}
```


temp = value;
value = temp + 1;
return temp;



Not so good...



Is this problem inherent?



If we could only glue reads and writes together...

Challenge

```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        temp = value;  
        value = temp + 1;  
        return temp;  
    }  
}
```


Challenge

```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        temp = value;  
        value = temp + 1;  
        return temp;  
    }  
}
```

**Make these steps
atomic (indivisible)**

Hardware Solution

```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        temp = value;  
        value = temp + 1;  
        return temp;  
    }  
}
```



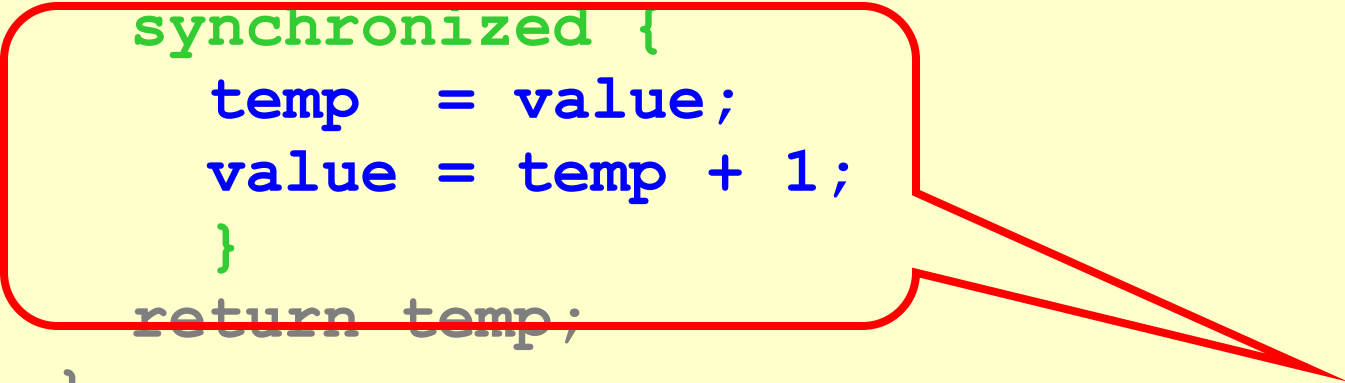
**ReadModifyWrite()
instruction**

An Aside: Java™

```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        synchronized {  
            temp = value;  
            value = temp + 1;  
        }  
        return temp;  
    }  
}
```

An Aside: Java™

```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        synchronized {  
            temp = value;  
            value = temp + 1;  
        }  
        return temp;  
    }  
}
```



Synchronized block

An Aside: Java™

```
public class Counter {  
    private long value;
```

```
    public long getAndIncrement() {  
        synchronized {
```

```
            temp = value;  
            value = temp + 1;
```

```
        }  
        return temp;
```

```
    }
```

```
}
```

Mutual Exclusion

